

Reinforcement Learning Training 2025

Model-Free Approach

Motivation

Recall in policy iteration

$$v_{k+1}(s) \leftarrow \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_k(s')]$$

- To make this work, we need to know the model dynamics or $p(s', r | s, a)$.
- However, we do not know p .
- Instead, we will resort to *sampling*.
 - Collecting experience by following some policy in the real world or running the agent through a policy in simulation.

Model-Free Learning

- Monte Carlo (MC) methods
- Temporal difference (TD) methods

Monte Carlo

- We use the law of large numbers (LLN) from statistics.
 - Average of samples is a good estimate for the actual unknown quantity.
 - This estimate becomes better and better as the number of trials of the experiment (samples) increases.

Monte Carlo


- Recall that we want to calculate

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$

- We let the agent start from this state $S_t = s$, follow the policy π to take actions, and keep doing so until termination.
 - We call one round of actions an **episode**.
- We record the total sum of rewards for each episode.
- We average the rewards to get an estimate of $v_{\pi}(s)$ for the policy π .

MC methods replace expected returns with the average of sample returns.

Worked Example



Solution v

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

let $\gamma=1$

$$V(C) = \boxed{0} \text{ (Terminal)}$$

$$\begin{aligned} V(B) &= 1 \left[1 \times [2 + 1(0)] \right] \\ &= \boxed{2} \end{aligned}$$

$$\begin{aligned} V(A) &= \frac{1}{3} [1 \times (0 + 1(2))] \\ &\quad + \frac{2}{3} [1 \times (1 + 1(0))] \\ &= \frac{1}{3}(2) + \frac{2}{3}(1) = \boxed{4/3} \end{aligned}$$

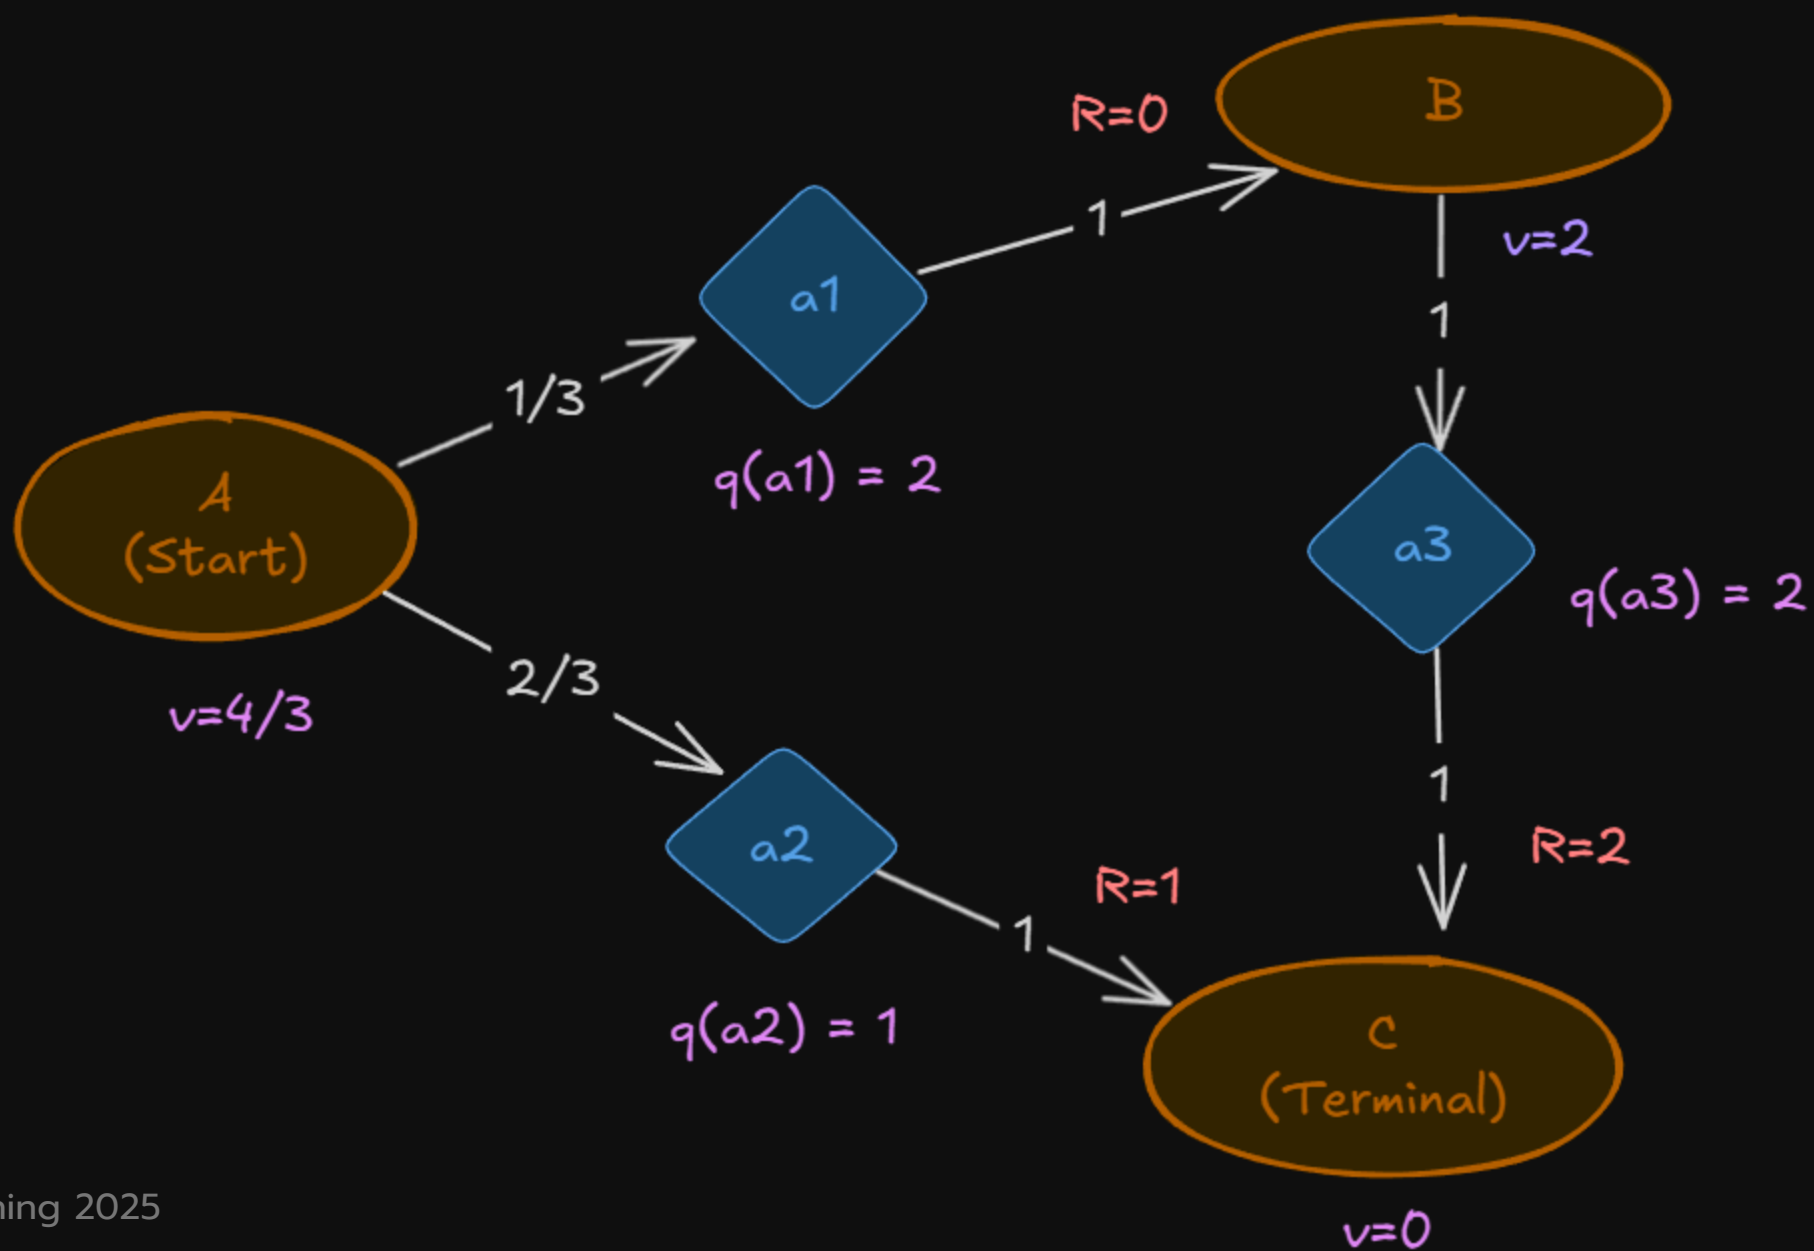
Solution q

$$q_{\pi}(s, a) = \sum_{s', r} p(s', r | s, a) \left[r + \gamma \sum_{a'} \pi(a' | s') q(s', a') \right]$$

$$\begin{aligned} q(a_3) &= 1 \left[2 + 1 \cancel{\sum(\dots)} \right] \\ &= \boxed{2} \end{aligned}$$

$$\begin{aligned} q(a_2) &= 1 \left[1 + 1 \cancel{\sum(\dots)} \right] \\ &= \boxed{1} \end{aligned}$$

$$\begin{aligned} q(a_1) &= 1 \left[0 + 1 (1(2)) \right] \\ &= \boxed{2} \end{aligned}$$



Estimate $v(A)$

- We simulate many episodes.

Episode	Path	Reward from A
1	$A \rightarrow C$	$G_1 = 1$
2	$A \rightarrow B \rightarrow C$	$G_2 = 0 + 2 = 2$
3	$A \rightarrow B \rightarrow C$	$G_3 = 0 + 2 = 2$
4	$A \rightarrow C$	$G_4 = 1$
...	...	G_n

Results

Monte Carlo estimates the value function $v(A)$ as the average return observed after visiting A.

$$v(A) = \frac{G_1 + G_2 + G_3 + G_4 + \dots}{n} = \frac{1 + 2 + 2 + 1 + \dots}{n} \rightarrow \frac{4}{3}$$

Online method

- Instead of averaging all the returns at the end (the sample mean), we can use the incremental (update) method to estimate $v(A)$ as each new return is observed.
- This is also called the "sample-average" update and is given by:

$$v_{n+1} = v_n + \frac{1}{n}(G_n - v_n)$$

Estimate $q(a_1)$ and $q(a_2)$

Episode	Path	Actions at A	Reward from Action at A
1	$A \rightarrow C$	a_2	$G_1 = 1$
2	$A \rightarrow B \rightarrow C$	a_1	$G_2 = 0 + 2$
3	$A \rightarrow B \rightarrow C$	a_1	$G_3 = 0 + 2$
4	$A \rightarrow C$	a_2	$G_4 = 1$
...	G_n

Estimate $q(a_1)$ and $q(a_2)$

$$q(a_1) = \frac{G_2 + G_3 + \dots}{n} = \frac{2 + 2 + \dots}{n} \rightarrow 2$$

$$q(a_2) = \frac{G_1 + G_4 + \dots}{n} = \frac{1 + 1 + \dots}{n} \rightarrow 1$$

Online update

$$q_{n+1} = q_n + \frac{1}{n}(G_n - q_n)$$

Comparing estimations of v and q

- Notice that the calculation of v and q is the same.
- This is because both functions are fundamentally estimates of an expected value—just over different types of returns.
 - $v(s)$ - average over all times you start at s
 - $q(s, a)$ - average over all times you start at s and pick a